

```
[28]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from wordcloud import WordCloud, STOPWORDS

In [2]: pip install wordcloud

Note: you may need to restart the kernel to use updated packages.

Collecting wordcloud
  Downloading wordcloud-1.9.3-cp39-cp39-win_and4.whl (389 kB)
Requirement already satisfied: matplotlib in c:\users\nitish singh\anaconda3\lib\site-packages (from wordcloud) (3.5.1)
Requirement already satisfied: pillow in c:\users\nitish singh\anaconda3\lib\site-packages (from wordcloud) (9.0.1)
Requirement already satisfied: numpy>=1.6.1 in c:\users\nitish singh\anaconda3\lib\site-packages (from wordcloud) (1.26.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (4.25.0)
Requirement already satisfied: python-dateutil<=2.7 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (2.8.2)
Requirement already satisfied: kiwisolver<=1.0.1 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (1.3.2)
Requirement already satisfied: packaging<=20.0 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (21.3)
Requirement already satisfied: pyparsing<=2.1.1 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (3.0.4)
Requirement already satisfied: cycler<=0.10 in c:\users\nitish singh\anaconda3\lib\site-packages (from matplotlib-wordcloud) (0.11.0)
Requirement already satisfied: six<=1.15 in c:\users\nitish singh\anaconda3\lib\site-packages (from python-dateutil<=2.7;matplotlib-wordcloud) (1.16.0)
Successfully installed wordcloud-1.9.3

In [29]: df = pd.read_csv("C:\Users\NITISH SINGH\Downloads\Reviews_data_dump\reviews_tawa\data.csv")
df.head()

Out[29]:
bound method NDFrame.head of
   Reviewer_Name  Reviewer_Rating  Review_Title
0      Sumit Kumar                5.0      Wonderful
1  BHARAT GALAGALI                5.0  Mind-blowing purchase
2  Paramjeet Singh                5.0      Awesome
3      Virendra Kumar                5.0      Great product
4      Jyoti solanki                5.0  Best in the market!
...
2526  Sandeep Mohapatra                3.0      Fair
2527  Vakul Rana                    4.0  Good quality product
2528  Sonraj Dhungana                4.0      Really Nice
2529  Veeranna N C                  3.0      Good
2530  Flipkart Customer                3.0      Just okay

...
Review_Text
0  I think in this price category it's best dosa ...
1  perfect tawa for Dosa. READ MORE
2  Excellent tawa. Made Paneer Tikka on first day...
3  Nice @ productREAD MORE
4  Delivery man is also good... READ MORE
...
2526  No Handel includedREAD MORE
2527  niceREAD MORE
2528  Everything is good in this price but handle is...
2529  Tawa OK .but Handel not regid.poor handle.READ...
2530  niceREAD MORE

...
Place_of_Review  Date_of_Review  Up_Votes
0      Certified Buyer, Lakhisarai      Sumit Kumar                211
1      Certified Buyer, Hunsur      BHARAT GALAGALI                107
2      Certified Buyer, Rampura Phul      Paramjeet Singh                59
3      Certified Buyer, Chengalpattu District      Virendra Kumar                77
4      Certified Buyer, Mumbai      Jyoti solanki                53
...
2526  Certified Buyer, Puri      Sandeep Mohapatra                0
2527  Certified Buyer, Chandigarh      Vakul Rana                0
2528  Certified Buyer, Bengaluru      Sonraj Dhungana                0
2529  Certified Buyer, Bengaluru      Veeranna N C                0
2530  Certified Buyer, Dharmapuri District      Flipkart Customer                2

...
Down_Votes
0      39
1      17
2      8
3      12
4      7
...
2526  0
2527  0
2528  0
2529  0
2530  1
dtypes: float64(1), int64(2), object(5)
memory usage: 158.3+ KB

In [31]:

In [31]:

In [31]:

In [32]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2531 entries, 0 to 2530
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
 #   Reviewer_Name        2531 non-null   object
 1  Reviewer_Rating       2285 non-null   float64
 2  Review_Title          2531 non-null   object
 3  Review_Text           2531 non-null   object
 4  Place_of_Review       2531 non-null   object
 5  Date_of_Review        2531 non-null   object
 6  Up_Votes              2531 non-null   int64
 7  Down_Votes            2531 non-null   int64
dtypes: float64(1), int64(2), object(5)
memory usage: 158.3+ KB

In [31]:

In [31]:

In [32]: df.isnull().sum()

Out[32]:
Reviewer_Name      0
Reviewer_Rating    246
Review_Title       0
Review_Text        0
Place_of_Review   0
Date_of_Review    0
Up_Votes          0
Down_Votes        0
dtype: int64

In [33]: # Assuming df is your DataFrame containing the data
# Fill missing values in 'Reviewer_Rating' column with mode
mode_val = df['Reviewer_Rating'].mode()[0]
df['Reviewer_Rating'].fillna(mode_val, inplace=True)

print(df)

Reviewer_Name  Reviewer_Rating  Review_Title \
0      Sumit Kumar                5.0      Wonderful
1  BHARAT GALAGALI                5.0  Mind-blowing purchase
2  Paramjeet Singh                5.0      Awesome
3      Virendra Kumar                5.0      Great product
4      Jyoti solanki                5.0  Best in the market!
...
2526  Sandeep Mohapatra                3.0      Fair
2527  Vakul Rana                    4.0  Good quality product
2528  Sonraj Dhungana                4.0      Really Nice
2529  Veeranna N C                  3.0      Good
2530  Flipkart Customer                3.0      Just okay

...
Review_Text
0  I think in this price category it's best dosa ...
1  perfect tawa for Dosa. READ MORE
2  Excellent tawa. Made Paneer Tikka on first day...
3  Nice @ productREAD MORE
4  Delivery man is also good... READ MORE
...
2526  No Handel includedREAD MORE
2527  niceREAD MORE
2528  Everything is good in this price but handle is...
2529  Tawa OK .but Handel not regid.poor handle.READ...
2530  niceREAD MORE

...
Place_of_Review  Date_of_Review  Up_Votes
0      Certified Buyer, Lakhisarai      Sumit Kumar                211
1      Certified Buyer, Hunsur      BHARAT GALAGALI                107
2      Certified Buyer, Rampura Phul      Paramjeet Singh                59
3      Certified Buyer, Chengalpattu District      Virendra Kumar                77
4      Certified Buyer, Mumbai      Jyoti solanki                53
...
2526  Certified Buyer, Puri      Sandeep Mohapatra                0
2527  Certified Buyer, Chandigarh      Vakul Rana                0
2528  Certified Buyer, Bengaluru      Sonraj Dhungana                0
2529  Certified Buyer, Bengaluru      Veeranna N C                0
2530  Certified Buyer, Dharmapuri District      Flipkart Customer                2

...
Down_Votes
0      39
1      17
2      8
3      12
4      7
...
2526  0
2527  0
2528  0
2529  0
2530  1
dtypes: float64(1), int64(2), object(5)
memory usage: 158.3+ KB

In [34]: df.isna().sum()

Out[34]:
Reviewer_Name      0
Reviewer_Rating    0
Review_Title       0
Review_Text        0
Place_of_Review   0
Date_of_Review    0
Up_Votes          0
Down_Votes        0
dtype: int64

In [92]: df.isnull().sum()

Out[92]:
Reviewer_Name      0
Reviewer_Rating    0
Review_Title       0
Review_Text        0
Place_of_Review   0
Date_of_Review    0
Up_Votes          0
Down_Votes        0
dtype: int64

In [35]: df.drop_duplicates(inplace=True)

In [36]: import string

df['Review_Title'] = df['Review_Title'].apply(lambda x: x.lower())
df['Review_Title'] = df['Review_Title'].apply(lambda x: x.translate(str.maketrans('', '', string.punctuation)))

In [37]: df['Review_Title']

Out[37]:
0      wonderful
1  mindblowing purchase
2      awesome
3      great product
4      best in the market
...
2526  fair
2527  good quality product
2528  really nice
2529  good
2530  just okay
Name: Review_Title, Length: 2471, dtype: object

In [38]: sns.histplot(df['Reviewer_Rating'])
plt.title('Distribution of Reviewer Rating')
plt.show()

Distribution of Reviewer Rating

In [39]: Up_Votes_counts = df['Up_Votes'].value_counts()
print(Up_Votes_counts)

0      2345
1      20
2      20
3      12
4      9
5      4
6      3
7      2
8      2
9      2
10     2
11     2
12     1
13     1
14     1
15     1
16     1
17     1
18     1
19     1
20     1
21     1
22     1
23     1
24     1
25     1
26     1
27     1
28     1
29     1
30     1
31     1
32     1
33     1
34     1
35     1
36     1
37     1
38     1
39     1
40     1
41     1
42     1
43     1
44     1
45     1
46     1
47     1
48     1
49     1
50     1
51     1
52     1
53     1
54     1
55     1
56     1
57     1
58     1
59     1
60     1
61     1
62     1
63     1
64     1
Name: Up_Votes, dtype: int64

In [40]: Down_Votes_counts = df['Down_Votes'].value_counts()
print(Down_Votes_counts)

0      2378
1      35
2      21
3      9
4      3
5      3
6      3
7      3
8      3
9      2
10     2
11     2
12     2
13     1
14     1
15     1
16     1
17     1
18     1
19     1
20     1
21     1
22     1
23     1
24     1
25     1
26     1
27     1
28     1
29     1
30     1
31     1
32     1
33     1
34     1
35     1
36     1
37     1
38     1
39     1
40     1
41     1
42     1
43     1
44     1
45     1
46     1
47     1
48     1
49     1
50     1
51     1
52     1
53     1
54     1
55     1
56     1
57     1
58     1
59     1
60     1
61     1
62     1
63     1
64     1
Name: Down_Votes, dtype: int64

In [41]: from sklearn.feature_extraction.text import CountVectorizer

# Assuming 'df' is your DataFrame containing text data
text_data = df['Review_Text']
vectorizer = CountVectorizer()
feature_matrix = vectorizer.fit_transform(text_data)
feature_names = vectorizer.get_feature_names_out()

In [42]: feature_names

Out[42]:
array(['@read', '3d', '100', ..., 'youread', 'yr', 'yumi'], dtype=object)

In [43]: # Assuming df is your DataFrame containing the data
# Assuming df.Review_Text contains the text data you want to transform

# Create an instance of CountVectorizer
count_vectorizer = CountVectorizer()

# Fit and transform the text data
data_features = count_vectorizer.fit_transform(df.Review_Text)

# Now you can use data_features for further analysis or modeling

In [46]: data_features

Out[46]:
<2471x1289 sparse matrix of type '<class 'numpy.int64''
with 12893 stored elements in Compressed Sparse Row format>

In [47]: data_features.shape

Out[47]:
(2471, 1289)

In [48]: data_features.getnnz()

Out[48]:
12993

In [49]: density = (data_features.getnnz() * 100) / (data_features.shape[0] * data_features.shape[1])
print("Density of the Matrix: ", density)

Density of the matrix: 0.37593754943958996

In [50]: feature_counts = df['Review_Text'].value_counts()
feature_counts

Out[50]:
258
niceREAD MORE
218
Good productREAD MORE
94
Nice productREAD MORE
79
Super READ MORE
64
...
Excellent product quantity, coating is also good, you can make healthy food on this tawa, & it has no indication base so it doesn't work in induc
on, overall good product. READ MORE 1
Satisfied @READ MORE
1
The product is good...liked it. .. Ok for daily useREAD MORE
1
Pan is good except for the handle which does not fix firmly. READ MORE 1
Tawa OK .but Handel not regid.poor handle. READ MORE 1
Name: Review_Text, Length: 1411, dtype: int64

In [64]: #Name the reviews with a function
import re
def cleanReviews(text):
    text = re.sub(r'@A-Za-z0-9_!@', '', text) # removes @mentions
    text = re.sub(r'!+', '', text) # removes exclamation '!' symbol
    text = re.sub(r'https?://\S+', '', text) # removes https
    text = re.sub(r'\n', ' ', text) # removes new line
    text = re.sub(r'\\n\\s+', ' ', text) # removes \\n\\s
    text = re.sub(r'\\.[!@|:|;|,|', ' ', text)
    text = re.sub(r'\\(|\\)|\\(|\\)', '', text)
    return text

df['cleanedReviews'] = df['Review_Text'].apply(cleanReviews) #apply cleanReviews function to the Reviews

df.head() #compares original reviews with cleaned Reviews

Out[64]:
Reviewer_Name  Reviewer_Rating  Review_Title  Review_Text  Place_of_Review  Date_of_Review  Up_Votes  Down_Votes  cleanedReviews
0      Sumit Kumar                5.0      wonderful  I think in this price category it's best dosa...  Certified Buyer, Lakhisarai      Sumit Kumar                211      39  I think in this price category it's best dosa...
1  BHARAT GALAGALI                5.0  mindblowing purchase  perfect tawa for Dosa. READ MORE  Certified Buyer, Hunsur      BHARAT GALAGALI                107      17  perfect tawa for Dosa READ MORE
2  Paramjeet Singh                5.0      awesome  Excellent tawa. Made Paneer Tikka on first day...  Certified Buyer, Rampura Phul      Paramjeet Singh                59      8  Excellent tawa Made Paneer Tikka on first day.
3  Virendra Kumar                5.0      great product  Nice @ productREAD MORE  Certified Buyer, Chengalpattu District      Virendra Kumar                77      12  Nice productREAD MORE
4      Jyoti solanki                5.0  best in the market  Delivery man is also good... READ MORE  Certified Buyer, Mumbai      Jyoti solanki                53      7  Delivery man is also good READ MORE

In [69]: df1=df[['cleanedReviews']]
df1

Out[69]:
cleanedReviews
0  I think in this price category it's best dosa L...
1  perfect tawa for Dosa READ MORE
2  Excellent tawa Made Paneer Tikka on first day...
3  Nice productREAD MORE
4  Delivery man is also good READ MORE
...
2526  No Handel includedREAD MORE
2527  niceREAD MORE
2528  Everything is good in this price but handle is...
2529  Tawa OK .but Handel not regid.poor handle READ...
2530  niceREAD MORE
2471 rows x 1 columns

In [70]: features = vectorizer.get_feature_names_out() # Replace with the variable that holds feature names
features_counts = np.sum(data_features.toarray(), axis=0)
features_counts_df = pd.DataFrame({'features': features, 'counts': features_counts})

.....
AttributeError                                Traceback (most recent call last)
...
1 features = vectorizer.get_feature_names_out() # Replace with the variable that holds feature names
...
2 features_counts = np.sum(data_features.toarray(), axis=0)
...
3 features_counts_df = pd.DataFrame({'features': features, 'counts': features_counts})
AttributeError: 'numpy.ndarray' object has no attribute 'toarray'

In [71]: plt.figure(figsize=(12, 5))
plt.hist(features_counts_df['counts'], bins=50, range=(0, 5000))
plt.xlabel("Frequency of Words")
plt.ylabel("Density")
plt.show()

Density

In [53]: count_of_single_occurrences = len(features_counts_df[features_counts_df['counts'] == 1])

Out[53]:
833

In [72]: count_vectorizer = CountVectorizer(max_features=10000)
feature_vector = count_vectorizer.fit_transform(df['cleanedReviews'])
features = count_vectorizer.get_feature_names_out()
data_features = feature_vector.toarray()
features_counts = np.sum(data_features, axis=0)
feature_counts = pd.DataFrame({'features': features, 'counts': features_counts})

In [73]: top_features_counts = feature_counts.sort_values('counts', ascending=False).head(15)

In [74]: top_features_counts

Out[74]:
features  counts
753  more      2483
515  good      626
997  read      458
2315  ve      425
942  rice      407
784  hot      380
952  products  372
520  goodness  369
514  s      265
476  hr      263
780  niceand  160
420  k      160
979  quality  159
161  rx      158
60  and      156

In [58]: pip install nltk

Requirement already satisfied: nltk in c:\users\nitish singh\anaconda3\lib\site-packages (3.7)
Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: click in c:\users\nitish singh\anaconda3\lib\site-packages (from nltk) (8.0.4)
Requirement already satisfied: regex<=2021.8.3 in c:\users\nitish singh\anaconda3\lib\site-packages (from nltk) (2022.3.15)
Requirement already satisfied: tqdm in c:\users\nitish singh\anaconda3\lib\site-packages (from nltk) (4.64.0)
Requirement already satisfied: joblib in c:\users\nitish singh\anaconda3\lib\site-packages (from nltk) (1.1.0)
Requirement already satisfied: colorama in c:\users\nitish singh\anaconda3\lib\site-packages (from click-nltk) (0.4.4)

In [59]: import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
print(nltk.download('stopwords.words('english')'))

[nltk_data] Downloading package stopwords to c:\users\NITISH SINGH\anaconda3\downloads\nltk_data...
[nltk_data]   UNZIPING corpora\stopwords.zip...

In [75]: df['cleanedReviews'][:10]

Out[75]:
0  I think in this price category it's best dosa t...
1  perfect tawa for Dosa READ MORE
2  Excellent tawa Made Paneer Tikka on first day...
3  Nice productREAD MORE
4  Delivery man is also good READ MORE
5  Now nice product Thank you Flipkart READ MORE
6  Value for money product. Thanks Let me write...
7  Nice product induction compatible must buy REA...
8  Nice @ READ MORE
9  Nice an good condition recive READ MORE
Name: cleanedReviews, dtype: object

In [68]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['cleanedReviews'], df['Reviewer_Rating'], test_size=0.2, random_state=42)

vectorizer = CountVectorizer()
vectorizer.fit(X_train)
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Train a classification model (Random Forest)
model = RandomForestClassifier()
model.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy: ", accuracy)
print("Classification Report:\n", report)

Accuracy: 0.6666666666666666
Classification Report:
precision    recall  f1-score   support

   3.0         0.00         0.07         0.12         60
   4.0         0.00         0.00         0.00         118
   5.0         0.65         1.00         0.79         317

 macro avg   0.48         0.36         0.65         495
weighted avg 0.51         0.65         0.82         495

C:\Users\NITISH SINGH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-de
fined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))
C:\Users\NITISH SINGH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-de
fined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))
C:\Users\NITISH SINGH\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-de
fined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  warn_pr(average, modifier, msg_start, len(result))

In [81]: import seaborn as sns
from sklearn.metrics import confusion_matrix

# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Create a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()

Confusion Matrix

In [87]: from sklearn.ensemble import RandomForestClassifier

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['cleanedReviews'], df['Reviewer_Rating'], test_size=0.2, random_state=42)

vectorizer = CountVectorizer()
vectorizer.fit(X_train)
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Train a classification model (Random Forest)
model = RandomForestClassifier()
model.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy: ", accuracy)
print("Classification Report:\n", report)

Accuracy: 0.6666666666666666
Classification Report:
precision    recall  f1-score   support

   3.0         0.50         0.15         0.24         60
   4.0         0.59         0.08         0.13         118
   5.0         0.67         0.97         0.79         317

 macro avg   0.58         0.40         0.39         495
weighted avg 0.62         0.66         0.45         495

In [89]: from sklearn.linear_model import LogisticRegression

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['cleanedReviews'], df['Reviewer_Rating'], test_size=0.2, random_state=42)

vectorizer = CountVectorizer()
vectorizer.fit(X_train)
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Train a classification model (Logistic Regression)
model = LogisticRegression()
model.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy: ", accuracy)
print("Classification Report:\n", report)

Accuracy: 0.6767676767676768
Classification Report:
precision    recall  f1-score   support

   3.0         0.65         0.25         0.36         60
   4.0         0.61         0.09         0.16         118
   5.0         0.68         0.97         0.80         317

 macro avg   0.65         0.44         0.44         495
weighted avg 0.63         0.68         0.60         495

C:\Users\NITISH SINGH\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: LBFGS failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = check_optimize_result(

In [91]: from sklearn.ensemble import GradientBoostingClassifier

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df['cleanedReviews'], df['Reviewer_Rating'], test_size=0.2, random_state=42)

vectorizer = CountVectorizer()
vectorizer.fit(X_train)
X_train_vectorized = vectorizer.fit_transform(X_train)
X_test_vectorized = vectorizer.transform(X_test)

# Train a classification model (Gradient Boosting)
model = GradientBoostingClassifier()
model.fit(X_train_vectorized, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_vectorized)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print the evaluation metrics
print("Accuracy: ", accuracy)
print("Classification Report:\n", report)

Accuracy: 0.6646464646464646
Classification Report:
precision    recall  f1-score   support

   3.0         0.78         0.12         0.20         60
   4.0         0.86         0.05         0.10         118
   5.0         0.66         1.00         0.79         317

 macro avg   0.76         0.36         0.45         495
weighted avg 0.68         0.68         0.45         495

These results indicate the performance of each model on sentiment classification. The accuracy values are consistent across all the classifiers. However, it is important to note that the models performed poorly in predicting the neutral sentiment, achieving an F1-score of 0.0 for this class. This suggests that the models might struggle to accurately classify neutral sentiments.

Further analysis and experimentation may be required to improve the performance of the models on the neutral sentiment class. Additionally, considering other evaluation metrics and exploring alternative approaches could provide deeper insights into the model's effectiveness for sentiment analysis.

Conclusions:
The sentiment analysis project aimed to classify customer reviews into awesome, very good, and average sentiments. Using various classifiers such as the SVM, Random Forest, Logistic Regression, and Gradient Boosting, the models achieved an accuracy of approximately 67%. However, the models struggled to accurately predict neutral sentiments. Further improvements and analysis are recommended.
```